

Masterpraktikum Scientific Computing – High Performance Computing

SWE Project Presentation and Results

Gaurav Kukreja
Evangelos Drossos

21.01.2014

Optimizations Performed

- Instrumentation with Scalasca
- Vectorization with Intel Cilk
- OpenMP optimization
- MPI optimization

Instrumentation

- Instrumented code to get a better idea of bottleneck issues and points of improvement
- Analyzed performance after each optimization

Instrumentation

- Initial Analysis
 - Showed that most significant time was spent in *computeNumericalFluxes()*
 - Waiting time for MPI Communication was significant

Vectorization

- Vectorization using Intel Cilk
 - Better Readable Code
 - Trust the compiler
- Split computation of rows into chunks. Computed chunks in one iteration of *computeNetUpdates()*
- After testing we determined an optimal chunk length of 8.

Vectorization

- Though we expected a large improvement, we achieved only slight improvement.
 - The compiler was already vectorizing the code.
 - Performance gets worse for larger block sizes. For block size 1, nearly same performance.
 - Based on tests on login node, performance $\sim 7\%$ better for chunk length 8

OpenMP


- Existing code is parallelized using OpenMP.
- Fixed a few compilation issues, and minor modifications
 - `schedule(static) nowait`
- Fused loops performing the calculations for the vertical and horizontal edge updates
- No cache optimization performed.

OpenMP

- Slight Improvement in performance.
 - ~1.5% against existing code.
 - Near Linear Speedup until 8 threads.
 - For more threads, slack expected.


OpenMP Optimizations

1000 x 1000



OpenMP Optimizations

1000 x 1000



Hybrid Approach


- Benefits of using MPI with OpenMP
 - OpenMP uses shared memory, which is good for cores on one node which share L2 Cache.
 - MPI provides flexibility and control for sending data between threads.
 - Hybrid approach can run optimally on multiple nodes.
 - Each MPI Thread runs on a node, and creates OpenMP Threads to run on the cores.

MPI

- Existing MPI Code used Blocking `MPI_Sendrecv()`
- To make Communication asynchronous and overlap with computation
 - Used non-blocking `MPI_Isend()` and `MPI_Irecv()`
 - Called `computeNumericalFluxes_innerBlock()` to compute horizontal and vertical edge updates for inner block
 - Wait and ensure communication was complete.
 - Called `computeNumericalFluxes_borders()`.
- ~12% improvement in performance

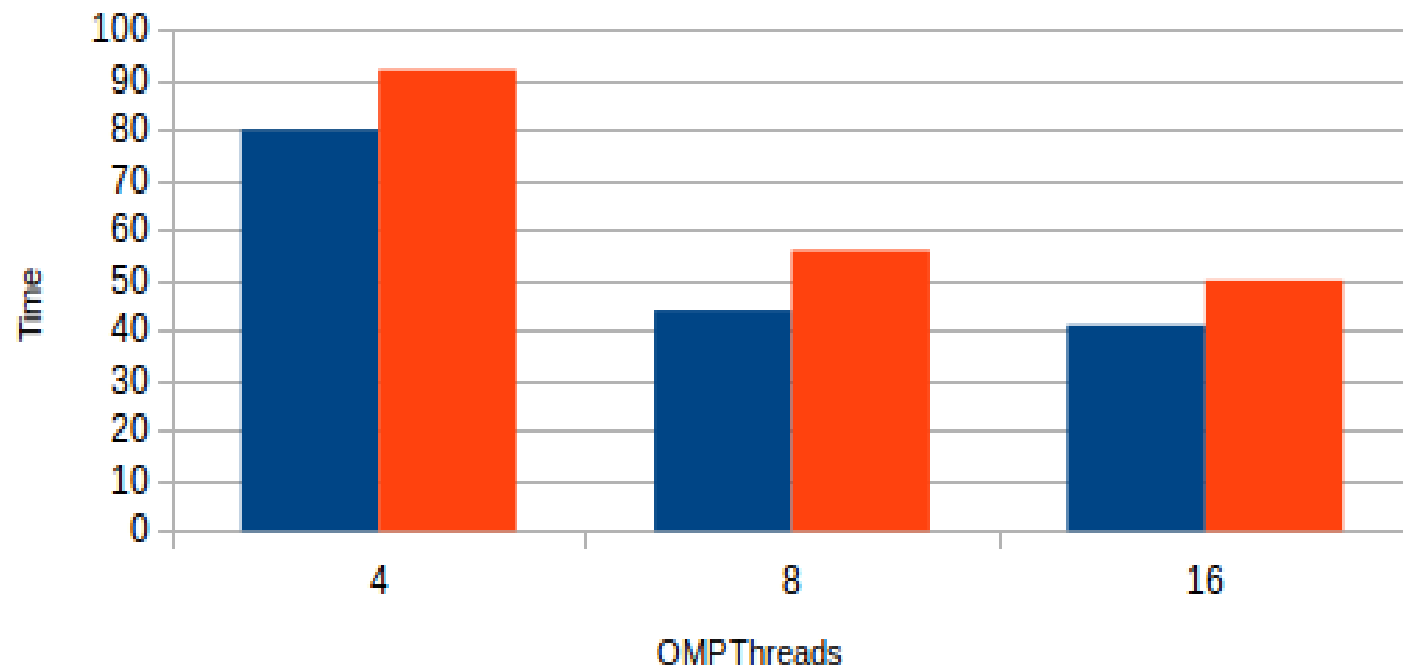
MPI Optimization

4 nodes with problem size 2000x2000



Improvement with MPI Optimization

4 nodes and 16 CPUs per node Problem Size 2000x2000



Hybrid Approach

- Worth noting
 - Hybrid code with problem size 2k x 2k on 4 nodes, and 8 CPUs takes slightly more time as OpenMP code with problem size 1k x 1k on 8 CPUs
 - The time difference is due to communication.
 - Overall improvement in performance ~10%